Group 3 week 3

Adrian, Cristina, Kira, and Mariana

Overview—Executive summary

- CIFAR-10 dataset
- Winning model:
  - Transfer Learning with VGG16 as base model
  - 86% accuracy on testing dataset

Dataset chosen

CIFAR-10 dataset:

60,000 images (32×32 pixels, RGB, 10 classes like airplane, car, cat, dog).

Why?

- Familiarity: we already used it in a lab.
- Size: This is a bigger dataset (more training data).
- Standardization: all images in the same shape and pixel scale [0, 255]

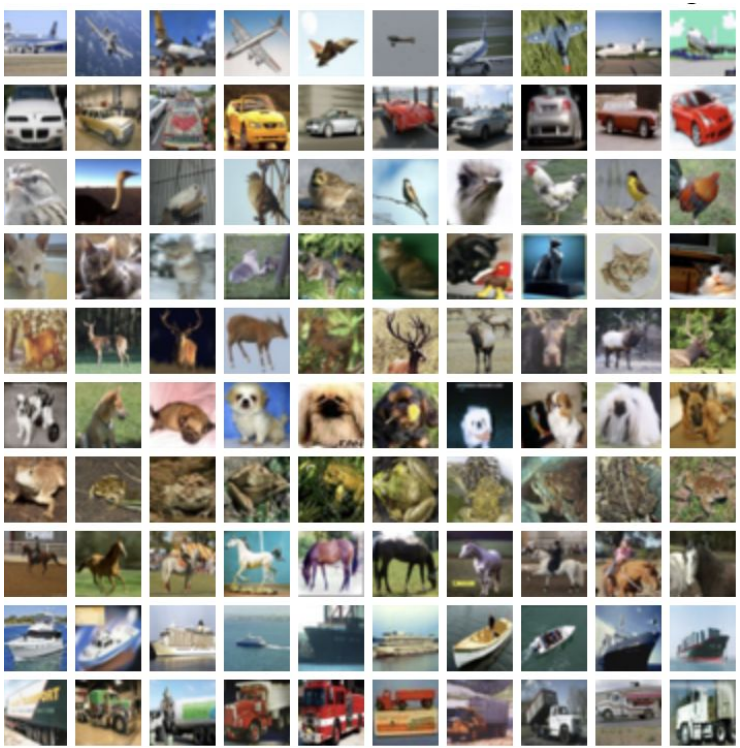

Data Preprocessing

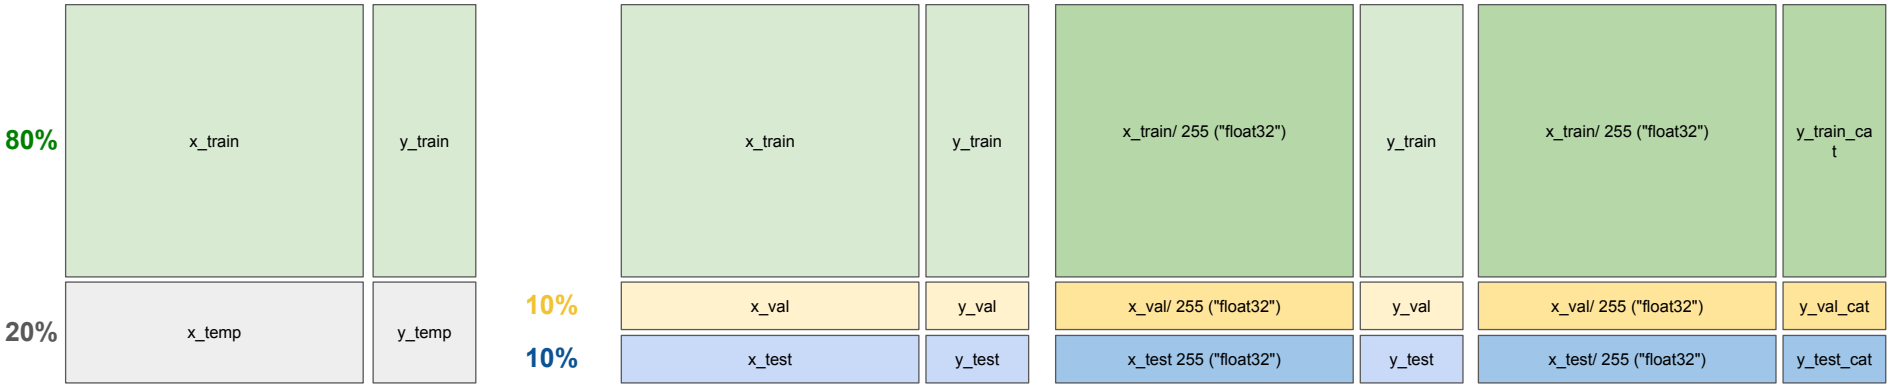

Original data from the dataset: already split into 2 groups.

Splitting data to keep some unused "real" examples

Scale images to the [0, 1]

convert classes into categories

Learning Process

	M1	M2	M3	M5	M17🏆
ITERATION BETS	Initial model	Added batch normalization	Increased number of epochs		Fixing the shape of the model (number of nodes, layers, etc).
PERFORMANCE change		4.08%	-1.22%	0.82%	11.88%
LEARNINGS		Good improvements	Performance dropped	Changed dropout	
Architecture	conv2d_5 (Conv2D)	conv2d_5 (Conv2D)	conv2d_5 (Conv2D)	conv2d_5 (Conv2D)	keras.Input(shape=input_sh
					#data_augmentation,
Hyperparamiters					
LR	9ms/s	9ms/s	9ms/s	9ms/s	
Optimizer	Adam	Adam	Adam	Adam	
epochs	15	15	50	100	
			Early Stopper	Early Stopper	
batch size	128	128	128	128	
Fit and eva					
Training time	3s 9ms/step	3s 9ms/step	3s 9ms/step		
Training loss					0.3396387994
Training accuracy					88.60%
Validation loss	0.8817585707	0.8089423776	1.171967149	0.7504937053	0.5229269266
Validatiion acc	69.56%	73.64%	72.42%	77.16%	81.44%
Test loss					0.5426046252
Test acc					81.44%

Initial model

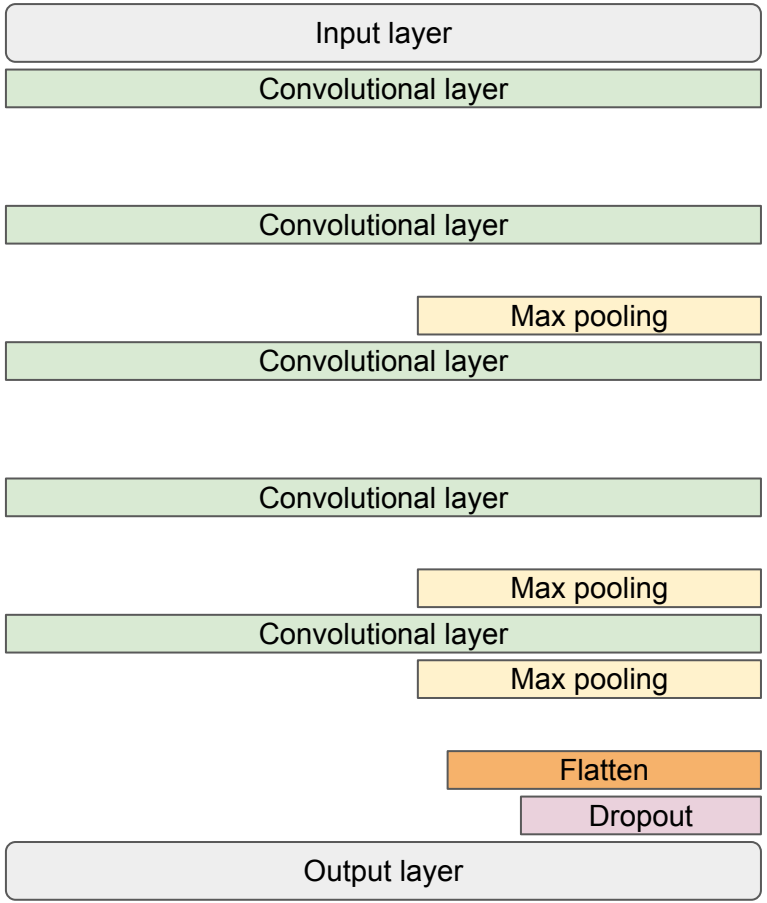

Validation loss: 0.8817585707
Validatiion acc: 69.56%

Tested 19 different models:

- 😐 More nodes in dense layers.
- 😊 Addition of batching normalization layers.
- 😊 Increasing number of epochs with an early stopper.
- 😐 Data augmentation.
- 😊 Additional dropout layers.
- 😊 Decreasing the dropout.
- 😐 Transfer learning did not work… at the beginning.

Best in-house model

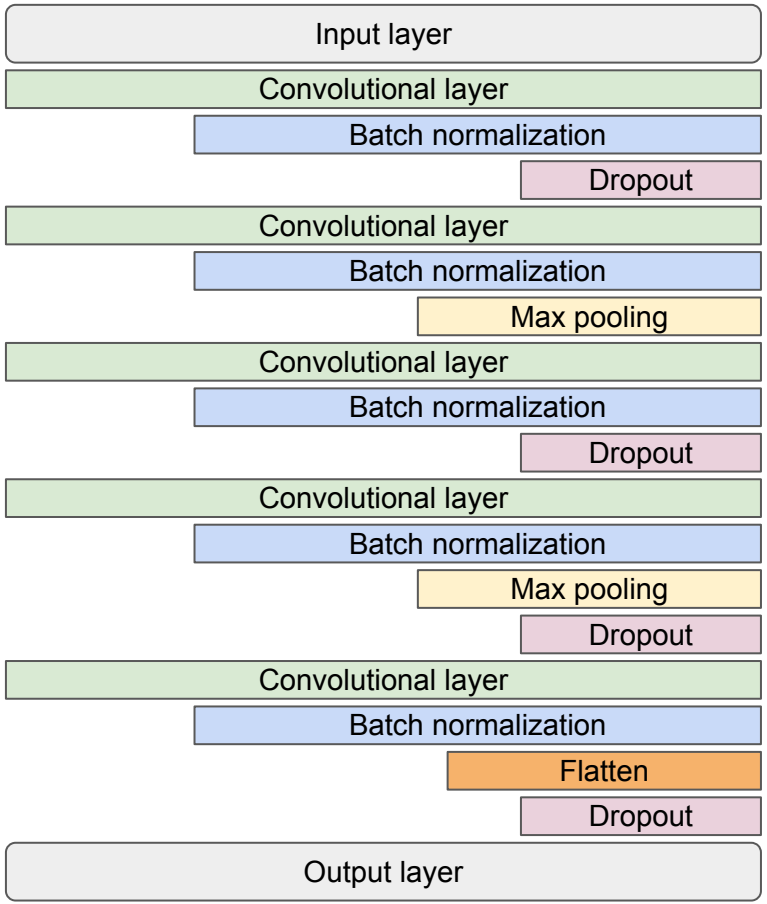

- **Avoid overfitting** (it delivers similar performance for the different dataset evaluations) by using dropout.
- Data augmentation did not improve our accuracy; our theory is that smaller-sized images can distort features so much that the model sees noise instead of useful patterns.
- The main performance increments came from the **batch normalization** and increment in **epochs** (from 15 to 100 with early stopper).

```
Train loss: 0.3465016484260559
Train accuracy: 0.8811799883842468
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 5ms/step
Validation loss: 0.5359759330749512
Validation accuracy: 0.8148000240325928
157/157 ━━━━━━━━━━━━━━━━━━━━ 0s 2ms/step
Test loss: 0.5468851327896118
Test accuracy: 0.8109999895095825
```

Transfer learning

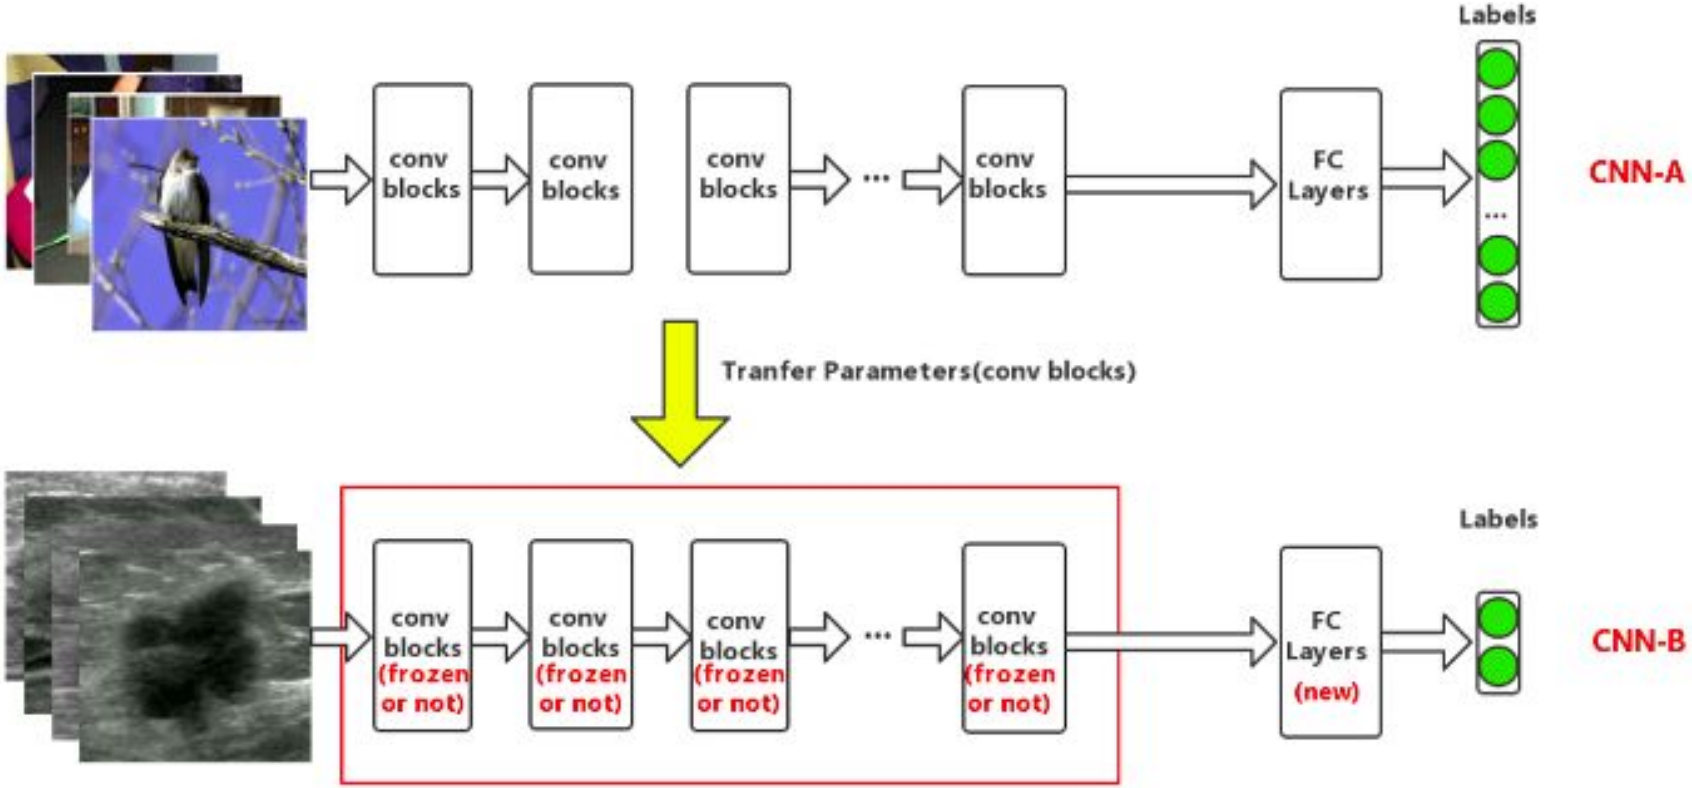

Winner model, aka Transfer Learning

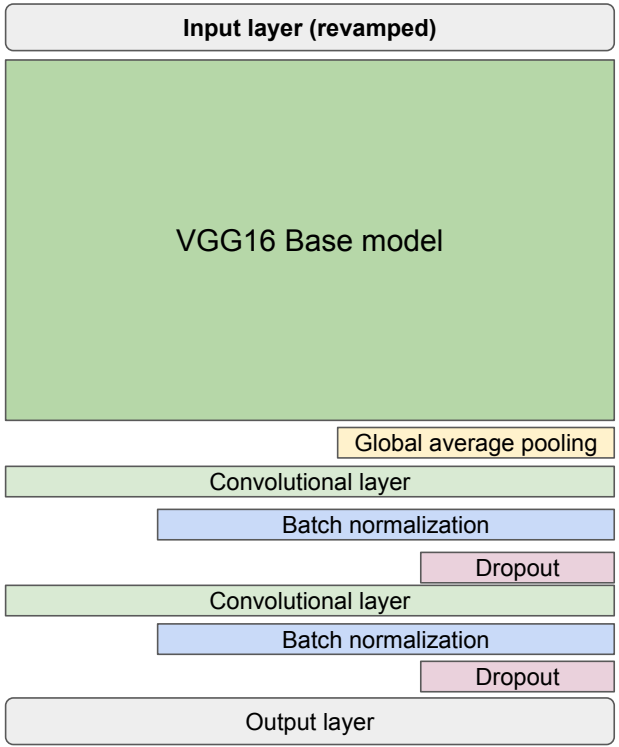

Validation loss: 0.34
Validation acc: 86.15%*
(*this model had only testing and training split.)

- **Preprocessing data** to fit base model
- **Global average pooling** (to convert the feature maps into vectors)
- Adding extra layers (**batch normalization** and **dropout**)
- Base model selection:
-

    1. MobileNetV2
    2. EfficientNETB0
    3. **VGG16**

Evaluation

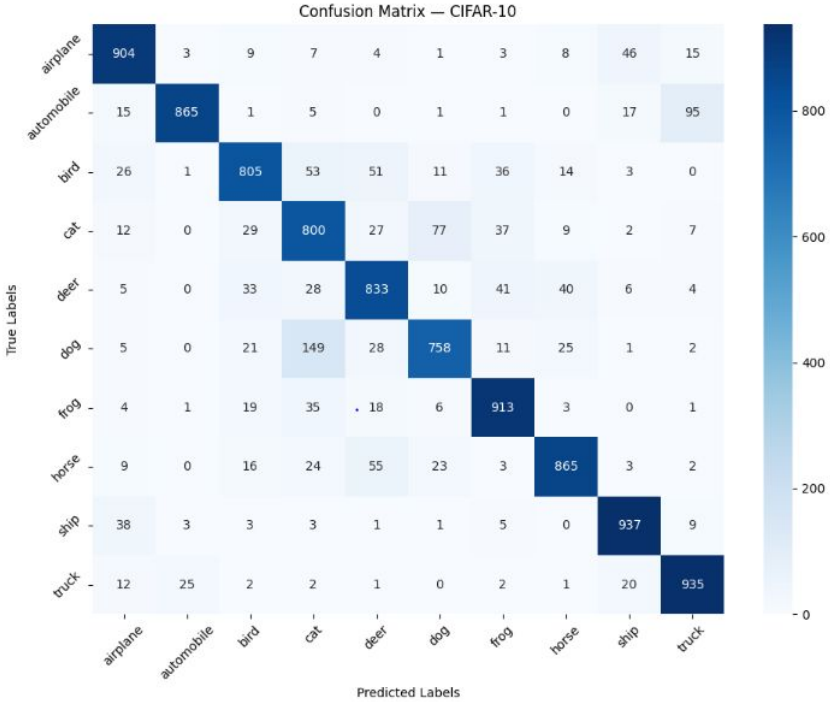

📊 Classification Report:

	precision	recall	f1-score	support
airplane	0.88	0.90	0.89	1000
automobile	0.96	0.86	0.91	1000
bird	0.86	0.81	0.83	1000
cat	0.72	0.80	0.76	1000
deer	0.82	0.83	0.83	1000
dog	0.85	0.76	0.80	1000
frog	0.87	0.91	0.89	1000
horse	0.90	0.86	0.88	1000
ship	0.91	0.94	0.92	1000
truck	0.87	0.94	0.90	1000
accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000

Learnings

- Iterative approach by changing one item at a time:
  - An increase of accuracy was achieved by adding batch normalization and dropout layers.
  - Avodied overfitting by adding dropout layers and early stopping.

- Selecting the dataset, we did not consider the challenges related to the sizing of images:
  - It affected the viability to use data augmentation.

- Selecting the correct base model in the data model: this is a trial and error to find the model that works better.

Thank you for your attention!

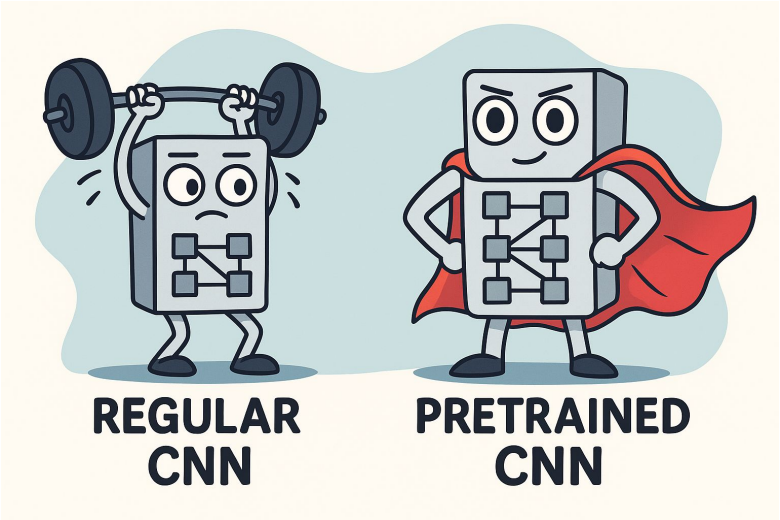

Mariana Borssatto, Cristina Insignares, Adrián A. H., Kira Redberg

Transfer learning process